

```

#!/usr/bin/env python
# coding: utf-8

# # Results Section

# In[1]:

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# In[2]:

df = pd.read_csv('ME202_Full_Data.csv') # doctest: +SKIP

# Summary of Data collected:

# In[3]:

df.info()

# In[4]:

df_data = df[["Title", "Page_Count", "Lowest_Price_Amazon"]]
df_data = df_data.dropna()
df_data = df_data.reset_index(drop=True)
df_data.Lowest_Price_Amazon = [x.strip('$') for x in
df_data.Lowest_Price_Amazon]
#df_data['Lowest_Price_Amazon'] =
df_data['Lowest_Price_Amazon'].to_numeric(col)
#df.Lowest_Price_Amazon = [x.strip('$') for x in
df.Lowest_Price_Amazon] #This should be faster
df_data[["Page_Count", "Lowest_Price_Amazon"]] =
df_data[["Page_Count", "Lowest_Price_Amazon"]].apply(pd.to_numeric)

# ### Data preparation
# Additional metadata was collected during the data collection
process. Only page count and lowest price data is used for regression.
These two variables were moved to a new dataframe. The book title was
also included in the dataframe for identification purposes.
#
# Summary of data in the modified dataframe:

```

```
# In[5]:
```

```
df_data.info()
```

```
# The statistics of the 2 variates collected were calculated:
```

```
# In[6]:
```

```
df_data.describe()
```

```
# The correlation coefficient of the 2 variables was calculated:
```

```
# In[69]:
```

```
#Correlation coefficient  
df_data["Page_Count"].corr(df_data["Lowest_Price_Amazon"])
```

```
# ### Plotting the Data
```

```
# Outliers can be identified when comparing relationships between two  
sets of data. A scatter plot was plotted to help identify any  
outliers.
```

```
#
```

```
# 2 outliers can be clearly seen.
```

```
# In[70]:
```

```
X = pages = df_data["Page_Count"]  
Y = price = df_data["Lowest_Price_Amazon"]
```

```
plt.figure("ro", figsize=(15,6))  
plt.title("Pages vs Price Plot", fontsize = 20)  
plt.plot(X, Y, 'ro')  
plt.xlabel('Number of Pages', fontsize=15)  
plt.ylabel('Price of book', fontsize = 15)
```

```
#plt.text(0,-200,"Figure 1: Pages vs Price Plot", fontsize = 15)  
plt.show()
```

```
# ### Fitting the regression line
```

```
# A regression line was then fitted on the dataset. The line of best  
fit was found with the normal equation.
```

```

#
# The Beta0 and Beta 1 obtained:

# In[9]:

X = np.array(X).reshape(-1, 1) # values converts it into a numpy
array
Y = np.array(Y).reshape(-1, 1) # -1 means that calculate the
dimension of rows, but have 1 column

from numpy.linalg import inv
def normal_equation(X, y):
    X_transpose = X.T
    theta = inv(X_transpose.dot(X)).dot(X_transpose).dot(y)

    # normal equation
    # theta_best = (X.T * X)^(-1) * X.T * y

    return theta # returns a list

X_b = np.c_[np.ones((len(df_data), 1)), X] # set bias term to 1 for
each sample
regression_params = normal_equation(X_b, Y)
print("beta0: ", regression_params[0])
print("beta1: ", regression_params[1])

#First one is intercept --> each dollar increase in price is
associated with an estimated increase in the number of pages of 1.3
pages
#Second one is slope --> maybe be interpreted as the predicted value
of the response when x = 0.
#This interpretation is only relevant when x = 0 is within the range
of the values of x from the sample

# In[10]:

def regression_predictor(regression_params, X):
    y_pred = regression_params[1] * X + regression_params[0]
    return y_pred

# The regression line was plotted with the dataset.

# In[73]:

```

```

Y_pred = regression_predictor(regression_params, X) # make
predictions
plt.figure("ro", figsize=(15,6))
plt.title("Regression Line for Pages to Price", fontsize = 20)
#plt.scatter(X, Y)
plt.plot(X, Y, 'ro')
plt.plot(X, Y_pred, color='blue')
plt.xlabel('Number of Pages', fontsize = 15)
plt.ylabel('Price of book', fontsize = 15)
#plt.text(0,-200,"Figure 2: Regression Line for Pages vs Price",
fontsize = 15);
plt.text(0,-200,"y = 0.14x + 51.22", fontsize = 15);

```

```

plt.show()

```

```

# ### Regression Line with outliers removed
# Since outliers can significantly affect the line of best fit for a
dataset, the outliers were removed from the data to explore the
effects they have on the regression line.

```

```

# In[12]:

```

```

#df_data.loc[df_data["Page_Count"] > 1200]

```

```

# In[13]:

```

```

#df_data.loc[df_data['Lowest_Price_Amazon'] > 800]

```

```

# In[14]:

```

```

#df_data.loc[df_data['Lowest_Price_Amazon'] > 800]
y_outlier_index = df_data[df_data['Lowest_Price_Amazon'] >
800].index.item()
#print(y_outlier_index)

```

```

# In[15]:

```

```

#df_data.loc[df_data["Page_Count"] > 1200]
x_outlier_index = df_data[df_data["Page_Count"] > 1200].index.item()

```

```
#print(x_outlier_index)
```

```
# In[58]:
```

```
#Outliers are saved in a separate dataframe for plotting later
outliers = pd.DataFrame(columns=df_data.columns)
row1 = df_data.iloc[x_outlier_index]
row2 = df_data.iloc[y_outlier_index]
outliers = outliers.append(row1, ignore_index=True)
outliers = outliers.append(row2, ignore_index=True)
```

```
outliers
```

```
# In[17]:
```

```
#Remove obvious outliers
df_data_no_outliers = df_data
df_data_no_outliers =
df_data_no_outliers.drop(df_data.index[y_outlier_index])
df_data_no_outliers =
df_data_no_outliers.drop(df_data.index[x_outlier_index])
df_data_no_outliers = df_data_no_outliers.reset_index(drop=True)
#df_data_no_outliers.head()
```

```
# The correlation coefficient was recalculated.
```

```
# In[18]:
```

```
#Correlation after removing outliers
df_data_no_outliers["Page_Count"].corr(df_data_no_outliers["Lowest_Price_Amazon"])
```

```
# The dataset without the outliers was replotted.
```

```
# In[74]:
```

```
X_no_outliers = pages = df_data_no_outliers["Page_Count"]
Y_no_outliers = price = df_data_no_outliers["Lowest_Price_Amazon"]

plt.figure("ro", figsize=(15,6))
plt.title("Pages vs Price Plot (Without outliers)", fontsize = 16)
plt.plot(X_no_outliers, Y_no_outliers, 'ro')
```

```

plt.xlabel('Number of Pages')
plt.ylabel('Price of book');

#plt.text(0,-60,"Figure 3: Pages vs Price Plot (Outliers removed)",
fontSize = 15)
plt.show()

# Another regression line was fitted on the dataset with outliers
removed.
#
# The new Beta0 and Beta 1 obtained:

# In[20]:

X_no_outliers = np.array(X_no_outliers).reshape(-1, 1) # values
converts it into a numpy array
Y_no_outliers = np.array(Y_no_outliers).reshape(-1, 1) # -1 means
that calculate the dimension of rows, but have 1 column

X_b_no = np.c_[np.ones((len(df_data_no_outliers), 1)), X_no_outliers]
# set bias term to 1 for each sample
regression_params_no_outliers = normal_equation(X_b_no, Y_no_outliers)

print("beta0: ", regression_params_no_outliers[0])
print("beta1: ", regression_params_no_outliers[1])

# The new regression line was plotted with the dataset with outliers
removed.

# In[78]:

Y_pred_no_outliers =
regression_predictor(regression_params_no_outliers, X_no_outliers) #
make predictions
plt.figure("ro", figsize=(15,6))
#plt.scatter(X, Y)
plt.title("Regression Line for Pages to Price (Outliers removed)",
fontSize = 16)
plt.plot(X_no_outliers, Y_no_outliers, 'ro')
plt.plot(X_no_outliers, Y_pred_no_outliers, color='blue')
plt.xlabel('Number of Pages')
plt.ylabel('Price of book')

#plt.text(0,-60,"Figure 4: Regression Line (Without outliers)",
fontSize = 15)

```

```
plt.text(0,-50,"y = 0.09x + 62.16", fontsize = 15);
```

```
plt.show()
```

```
# Both regression lines were plotted with the data and outliers.
```

```
# In[83]:
```

```
plt.figure("ro", figsize=(15,10))
#plt.scatter(X, Y)
plt.title("Regression Lines for Pages to Price", fontsize = 16)

X_outlier = outliers["Page_Count"]
Y_outlier = outliers["Lowest_Price_Amazon"]

plt.plot(X_no_outliers, Y_no_outliers, 'ro', label = "Dataset")
plt.plot(X_outlier, Y_outlier, "g^", label = "Outliers")
plt.plot(X, Y_pred, color='green', label="Regression line With
outliers")
plt.plot(X_no_outliers, Y_pred_no_outliers, color='blue', label=
"Regression line with outliers removed")
plt.legend()
plt.xlabel('Number of Pages')
plt.ylabel('Price of book')
#plt.text(0,-125,"Figure 5: Plot of both regression lines", fontsize =
15)
plt.show()
```

```
# The standard error for the two regression lines are calculated:
```

```
# In[23]:
```

```
def standard_error(x_input, y_input, y_pred):
    x_mean = np.sum(x_input) / len(x_input)

    x_difference_term = 0.0
    for i in range(len(x_input)):
        x_difference_term += (x_input[i]-x_mean)**2

    error_sum = 0.0
    for i in range(len(y_input)):
        error_sum = (y_input[i] - y_pred[i])**2

    top = error_sum / (len(x_input)-2)
```

```

    #print(top[0])
    #print(x_difference_term[0])
    return np.sqrt(top[0]/x_difference_term[0]) #standard error

```

In[24]:

```

SE = standard_error(X, Y, Y_pred)
SE_without_outliers = standard_error(X_no_outliers, Y_no_outliers,
Y_pred_no_outliers)

```

```

print("Standard Error with all data:", SE)
print("Standard Error with outliers removed", SE_without_outliers)

```

The 95% confidence interval for beta1 is determined

In[85]:

```

import scipy.stats as stats

dof = len(X_no_outliers) - 2

#95% confidence interval for beta1
t_n2 = stats.t.ppf(q=0.975, # Quantile to check
                    df=dof) # Degrees of freedom""""
t_n2 = float(t_n2)
print("t value:", t_n2)
#print(t_n2, type(t_n2))

beta1 = float(regression_params_no_outliers[1])
#print(type(beta1))

print(beta1)
print(t_n2*SE)

c_interval_min = beta1 - t_n2*SE
c_interval_max = beta1 + t_n2*SE

print("The Confidence Interval is:(",c_interval_min,
",",c_interval_max,")")

```

There is a 95% confidence that an increase in one page is associated with a mean increase of between 1 and 10 cents on the book's price.

If there is no relationship between the number of pages in a book and the book's price, beta1 will be equal to zero.


```

#
# A hypothesis test was conducted on beta1 to determine if there is a
# significant positive association. The values used to calculate beta1
# and standard error had the 2 outliers removed to avoid the results
# being skewed by the outliers.
#
# To setup the hypothesis test:
#
#      H0: B1 = 0
#      Ha: B1 != 0
#

# In[26]:

b1 = float(regression_params_no_outliers[1])
SE = standard_error(X_no_outliers, Y_no_outliers, Y_pred_no_outliers)

print("Beta1:", b1)
print("Standard error:", SE)

t = b1/SE
print("The t value can be calculated as beta1 divided by the standard
error.")
print("t value:", t)

dof = len(X_no_outliers) - 2
print("The degree of freedom is:", dof)

"""stats.t.ppf(q=0.025, # Quantile to check
               df=dof) # Degrees of freedom

stats.t.ppf(q=0.975, # Quantile to check
            df=dof) # Degrees of freedom"""

T_test = stats.t.cdf(x = t,      # T-test statistic
                    df= dof)

print("T test value:", T_test)

p_value = 2*(1-(T_test))
print("P value is:", p_value)
print()

if (p_value < 0.001):
    print("Since the p_value is less than 0.001,")
    print ("H0 is rejected.")
else:
    print("H0 is accepted. The p_value is greater than 0.001.")

```

```
# Based on the hypothesis test on b1, it can be concluded that there
is a significant positive association between the number of pages and
price of the book.
```

```
# In[ ]:
```

```
get_ipython().run_cell_magic('html', '', '<style>\ndiv.input {\n
display:none;\n}\n</style>')
```

```
# ## References:
```

```
#
```

```
# 1) http://reliawiki.org/index.php/  
Simple_Linear_Regression_Analysis#Fitted_Regression_Line
```

```
#
```

```
# 2)  
# http://hamelg.blogspot.com/2015/11/python-for-data-analysis-  
part-24.html
```